

Announcements

ITIS 6200 / 8200

- Complete Homework 0 if you haven't
- Teaching feedback welcome

Last Time: Introduction and Security Principles

ITIS 6200 / 8200

- Course Logistics
- Security Principles

Today's Cunning Plan: Cryptography

ITIS 6200 / 8200

- Introduction to Cryptography

- What is cryptography?
- Definitions of key properties
- IND-CPA security
- A brief history of cryptography
- Symmetric-key crypto
 - One-time pads

What is cryptography?

What is Cryptography?

ITIS 6200 / 8200

- **Older definition:** The study of secure communication over insecure channels
- **Newer definition:** Provide rigorous guarantees on the security of data and computation in the presence of an attacker
 - Not just *confidentiality* but also *integrity* and *authenticity* (we'll see these definitions today)
- Modern cryptography is heavily based on math
 - The math is hard
 - We won't cover the details. Wikipedia is a great resource
 - Some system / algorithms are provably secure

Definitions of Key Properties

Meet Alice, Bob, Eve, and Mallory

ITIS 6200 / 8200

- Alice and Bob: The main characters trying to send messages to each other over an insecure communication channel
- Eve: An **eavesdropper** who can read any data sent over the channel
- Mallory: A **manipulator** who can read and modify any data sent over the channel

Meet Alice, Bob, Eve, and Mallory

ITIS 6200 / 8200

- We often describe cryptographic problems using a common cast of characters
- One scenario:
 - Alice wants to send a message to Bob.
 - However, Eve is going to *eavesdrop* on the communication channel.
 - How does Alice send the message to Bob without Eve learning about the message?
- Another scenario:
 - Bob wants to send a message to Alice.
 - However, Mallory is going to *tamper* with the communication channel.
 - How does Bob send the message to Alice without Mallory changing the message?

Three Goals of Cryptography

ITIS 6200 / 8200

- In cryptography, there are three main properties that we want on our information
- **Confidentiality:** An adversary cannot *read* our messages.
- **Integrity:** An adversary cannot *change* our messages without being detected.
- **Authenticity:** I can prove that this message came from the person who claims to have written it.
 - Integrity and authenticity are closely related properties...
 - Before I can prove that a message came from a certain person, I have to prove that the message wasn't changed!

Keys

ITIS 6200 / 8200

- The most basic building block of any cryptographic scheme:
The **key**
- We use the key in our algorithms to secure messages
- Two models of keys:
 - **Symmetric key model:** Alice and Bob both know the value of the same secret key.
 - Example: One-time pad, AES encryption
 - **Asymmetric key model:** Everybody has two keys, a secret key and a public key.
 - Example: RSA encryption

Security Principle: Kerckhoff's Principle

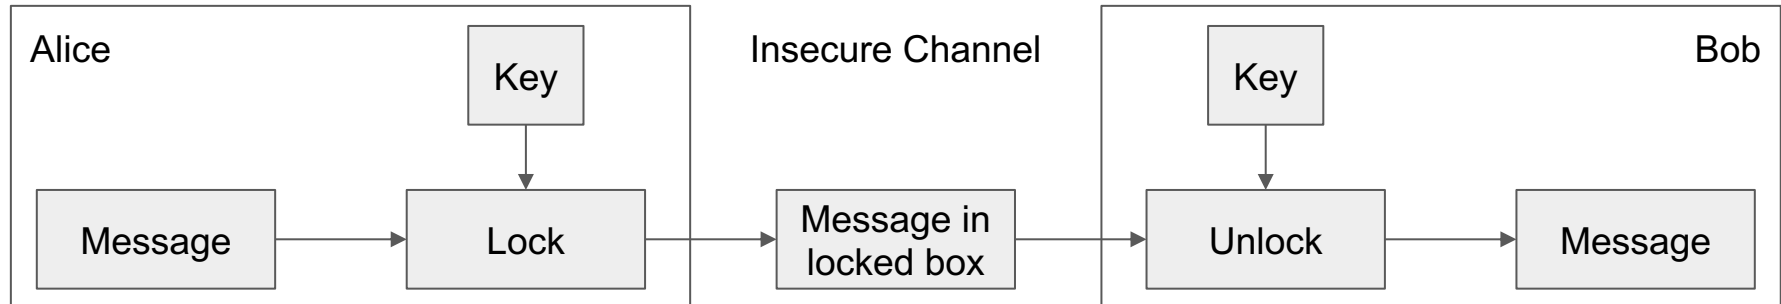
ITIS 6200 / 8200

- This principle is closely related to Shannon's Maxim
 - Don't use security through obscurity. Assume the attacker knows the system.
- Kerckhoff's principle says:
 - Cryptosystems should remain secure even when the attacker knows all internal details of the system
 - The key should be the only thing that must be kept secret
 - The system should be designed to make it easy to change keys that are leaked (or suspected to be leaked)
 - If your secrets are leaked, it is usually a lot easier to change the key than to replace every instance of the running software
- Our assumption: The attacker knows all the algorithms we use. The only information the attacker is missing is the secret key(s).

Confidentiality

ITIS 6200 / 8200

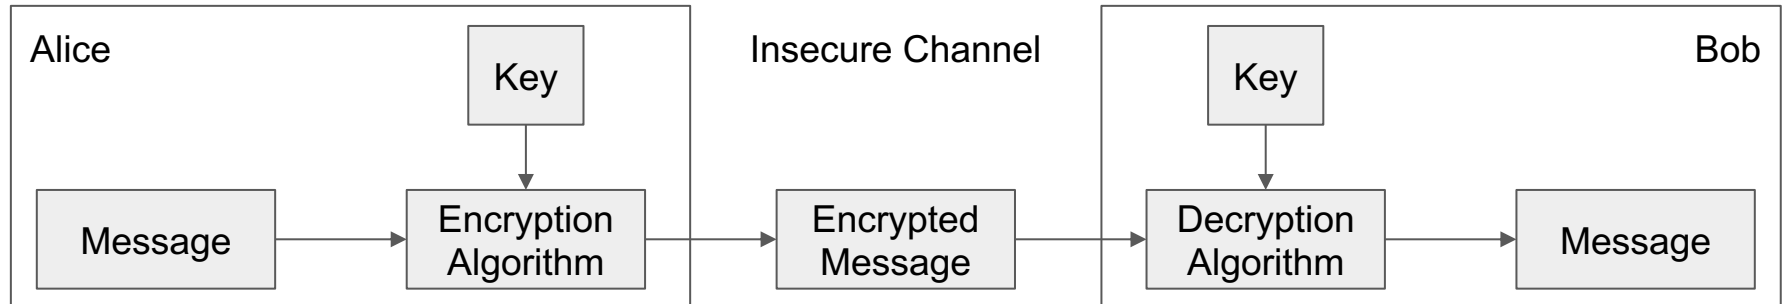
- **Confidentiality:** An adversary cannot *read* our messages.
- Analogy: Locking and unlocking the message
 - Alice uses the key to lock the message in a box
 - Alice sends the message (locked in the box) over the insecure channel
 - Eve sees the locked box, but cannot access the message without the key
 - Bob receives the message (locked in the box) and uses the key to unlock the message



Confidentiality

ITIS 6200 / 8200

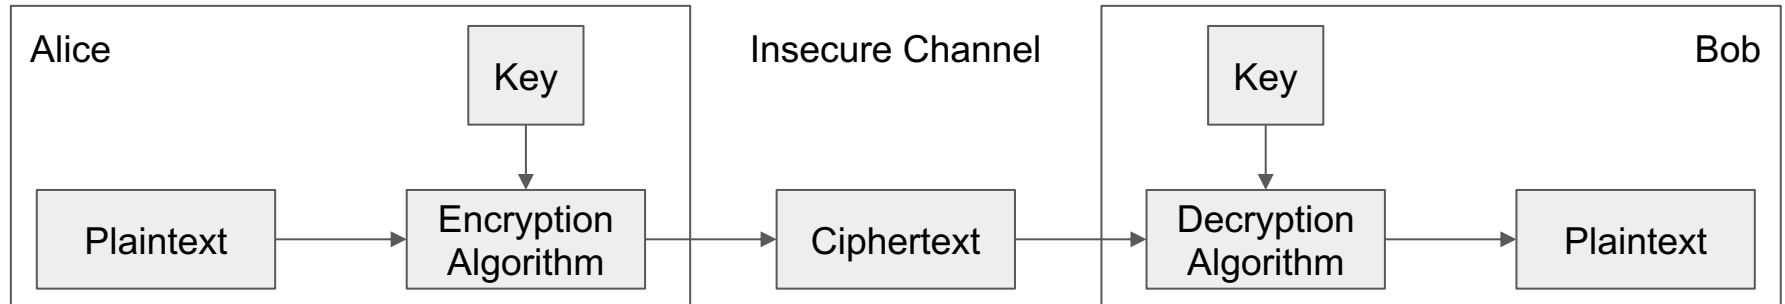
- Schemes provide confidentiality by **encrypting** messages
 - Alice uses the key to **encrypt** the message: Change the message into a scrambled form
 - Alice sends the encrypted message over the insecure channel
 - Eve sees the encrypted message, but cannot figure out the original message without the key
 - Bob receives the encrypted message and uses the key to **decrypt** the message back into its original form



Confidentiality

ITIS 6200 / 8200

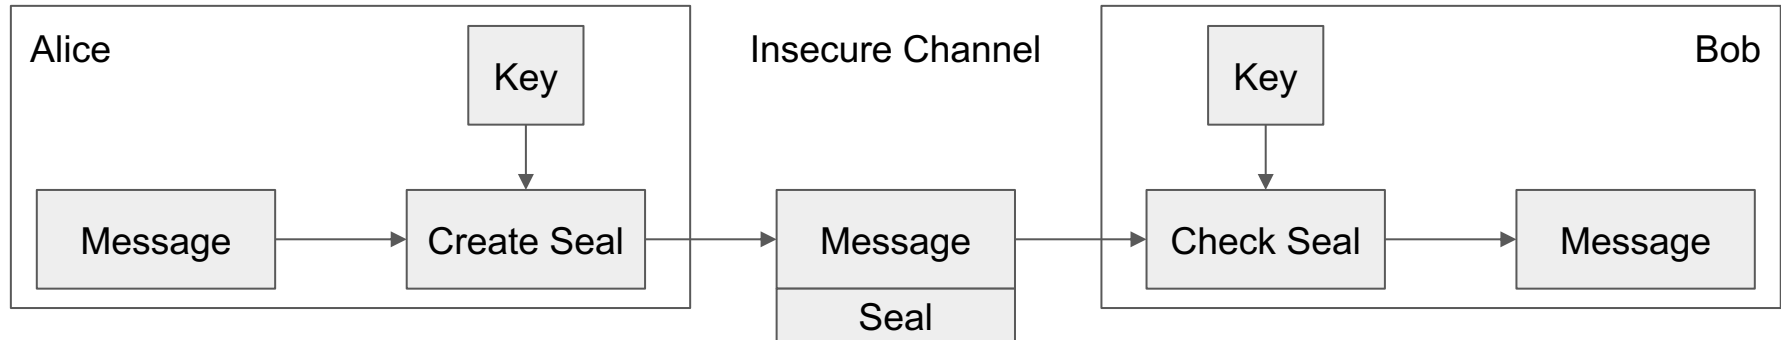
- **Plaintext:** The original message
- **Ciphertext:** The encrypted message



Integrity (and Authenticity)

ITIS 6200 / 8200

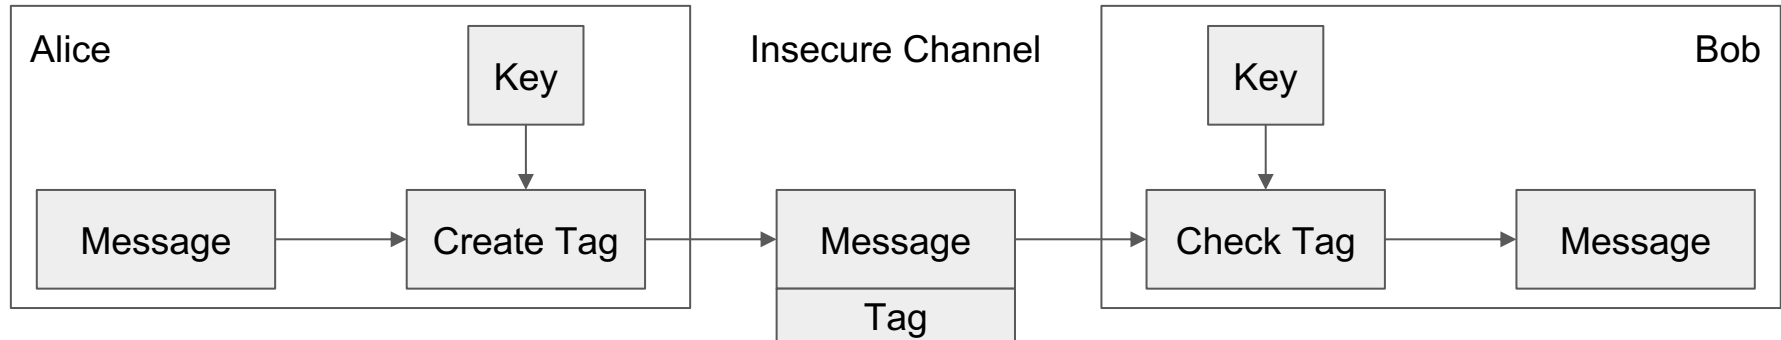
- **Integrity:** An adversary cannot *change* our messages without being detected.
- **Analogy: Adding a seal on the message**
 - Alice uses the key to add a special seal on the message (e.g. puts tape on the envelope)
 - Alice sends the message and the seal over the insecure channel
 - If Mallory tampers with the message, she'll break the seal (e.g. break the tape on the envelope)
 - Without the key, Mallory cannot create her own seal
 - Bob receives the message and the seal and checks that the seal has not been broken



Integrity (and Authenticity)

ITIS 6200 / 8200

- Schemes provide integrity by adding a **tag** or **signature** on messages
 - Alice uses the key to generate a special tag for the message
 - Alice sends the message and the tag over the insecure channel
 - If Mallory tampers with the message, the tag will no longer be valid
 - Bob receives the message and the tag and checks that the tag is still valid



Threat Models (Types of Cryptanalysis)

ITIS 6200 / 8200

- Ciphertext only
 - Attackers only have ciphertext, need to decrypt it
- Known-plaintext
 - Attackers know **some** ciphertext and its corresponding plaintext, needs to decrypt some other cyphertext
- Chosen-plaintext
 - Attackers can get **any** ciphertext they want encrypted, and need to decrypt some specific cyphertext
 - e.g., Eve can trick Alice into encrypting arbitrary messages of Eve's choice
- In this class, we often use the chosen plaintext attack (CPA) model

Cryptography Roadmap

ITIS 6200 / 8200

	Symmetric-key	Asymmetric-key
Confidentiality	<ul style="list-style-type: none">• One-time pads• Block ciphers with chaining modes (e.g. AES-CBC)• Stream ciphers	<ul style="list-style-type: none">• RSA encryption• ElGamal encryption
Integrity, Authentication	<ul style="list-style-type: none">• MACs (e.g. HMAC)	<ul style="list-style-type: none">• Digital signatures (e.g. RSA signatures)

- Hash functions
- Pseudorandom number generators
- Public key exchange (e.g. Diffie-Hellman)

- Key management (certificates)
- Password management

Symmetric-Key Encryption

Cryptography Roadmap

ITIS 6200 / 8200

	Symmetric-key	Asymmetric-key
Confidentiality	<ul style="list-style-type: none">• One-time pads• Block ciphers with chaining modes (e.g. AES-CBC)• Stream ciphers	<ul style="list-style-type: none">• RSA encryption• ElGamal encryption
Integrity, Authentication	<ul style="list-style-type: none">• MACs (e.g. HMAC)	<ul style="list-style-type: none">• Digital signatures (e.g. RSA signatures)

- Hash functions
- Pseudorandom number generators
- Public key exchange (e.g. Diffie-Hellman)

- Key management (certificates)
- Password management

Symmetric-Key Encryption

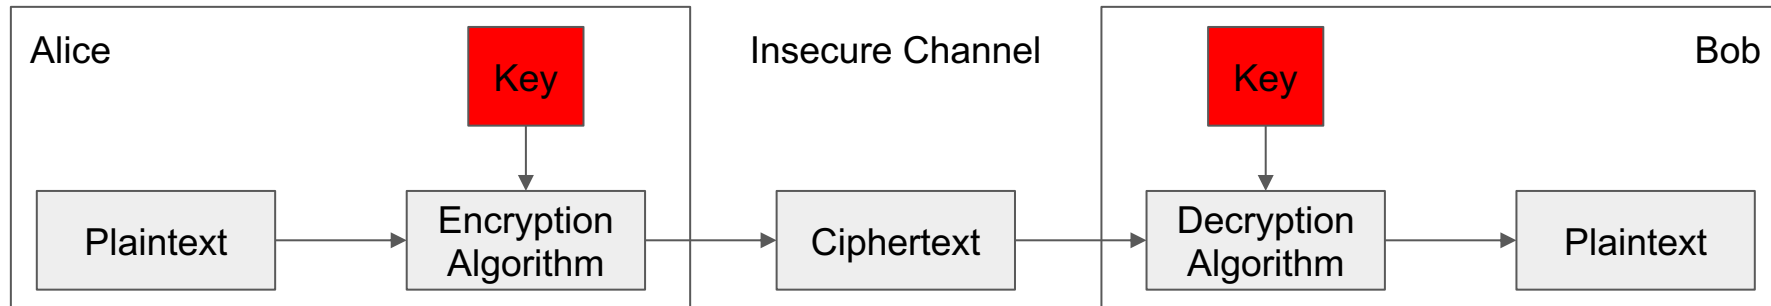
ITIS 6200 / 8200

- The next few schemes are symmetric-key encryption schemes
 - **Encryption schemes** aim to provide *confidentiality* (but not integrity or authentication)
 - **Symmetric-key** means Alice and Bob share the same secret key that the attacker doesn't know
 - Don't worry about how Alice and Bob share the key for now
- For modern schemes, we're going to assume that messages are *bitstrings*
 - **Bitstring**: A sequence of bits (0 or 1), e.g. 11010101001001010
 - Text, images, etc. can usually be converted into bitstrings before encryption, so bitstrings are a useful abstraction. After all, everything in a computer is just a sequence of bits!

Symmetric-Key Encryption: Definition

ITIS 6200 / 8200

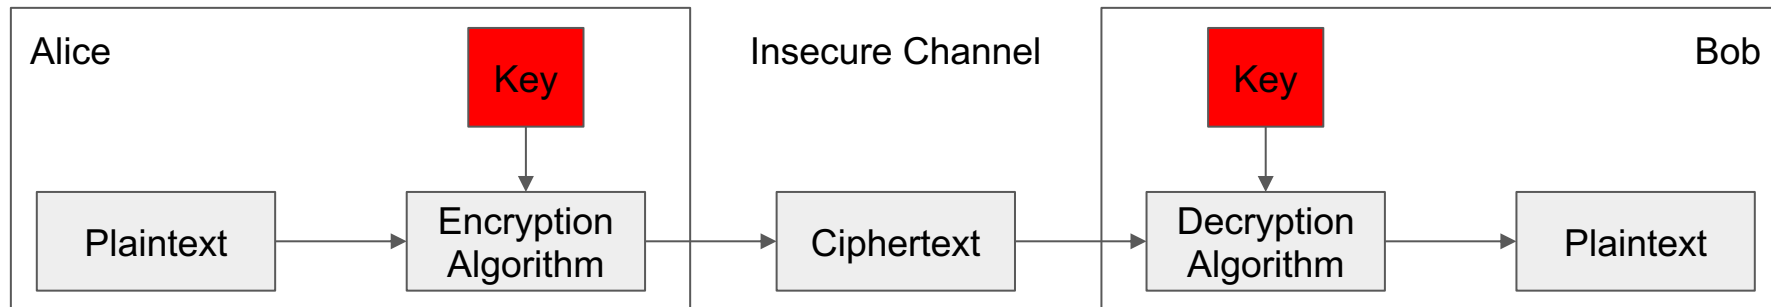
- A symmetric-key encryption scheme has three algorithms:
 - $\text{KeyGen}() \rightarrow K$: Generate a key K
 - $\text{Enc}(K, M) \rightarrow C$: Encrypt a **plaintext** M using the key K to produce **ciphertext** C
 - $\text{Dec}(K, C) \rightarrow M$: Decrypt a ciphertext C using the key K



Symmetric-Key Encryption: Definition

ITIS 6200 / 8200

- What properties do we want from a symmetric encryption scheme?
 - **Correctness:** Decrypting a ciphertext should result in the message that was originally encrypted
 - $\text{Dec}(K, \text{Enc}(K, M)) = M$ for all $K \leftarrow \text{KeyGen}()$ and M
 - **Efficiency:** Encryption/decryption algorithms should be fast: >1 Gbps on a standard computer
 - **Security:** Confidentiality



Defining Confidentiality

ITIS 6200 / 8200

- Recall our definition of confidentiality from earlier: “An adversary cannot read our messages”
 - This definition isn’t very specific
 - What if Eve can read the first half of Alice’s message, but not the second half?
 - What if Eve figures out that Alice’s message starts with “Dear Bob”?
 - This definition doesn’t account for prior knowledge
 - What if Eve already knew that Alice’s message ends in “Sincerely, Alice”?
 - What if Eve knows that Alice’s message is “BUY!” or “SELL” but doesn't know which?



Defining Confidentiality

ITIS 6200 / 8200

- A better definition of confidentiality: The ciphertext should not give the attacker *any **additional** information* about the plaintext.
- Let's design an experiment/security **game** to test our definition:
 - Eve chooses two messages M_0 and M_1 of the same length
 - Alice chooses one message at random M_b , encrypts it, and sends the ciphertext
 - Eve knows either M_0 or M_1 was sent, but doesn't know which
 - Eve reads the ciphertext and tries to guess which message was sent
 - If the probability that Eve correctly guesses which message was sent is $1/2$, then the encryption scheme is confidential
- Intuition
 - If the scheme is confidential, Eve can only guess with probability $1/2$, which is no different than if Eve hadn't sent the ciphertext at all
 - In other words: the ciphertext gave Eve no **additional** information about which plaintext was sent!

Defining Confidentiality: IND-CPA

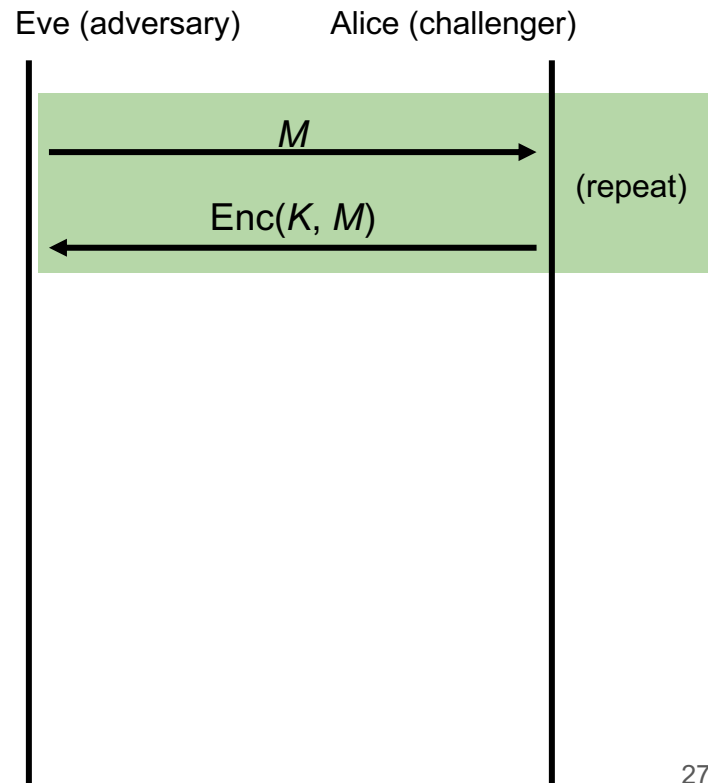
ITIS 6200 / 8200

- Recall our threat model: Eve can also perform a **chosen plaintext attack (CPA)**
 - It means, Eve can trick Alice into encrypting arbitrary messages of Eve's choice
 - We can adapt our experiment to account for this threat model
- A better definition of confidentiality: Even if Eve is able to trick Alice into encrypting messages, Eve can still only guess what message Alice sent with probability $1/2$.
 - This definition is called **IND-CPA** (indistinguishability under chosen plaintext attack)
- Cryptographic properties are often defined in terms of “games” that an adversary can either “win” or “lose”

Defining Confidentiality: IND-CPA

ITIS 6200 / 8200

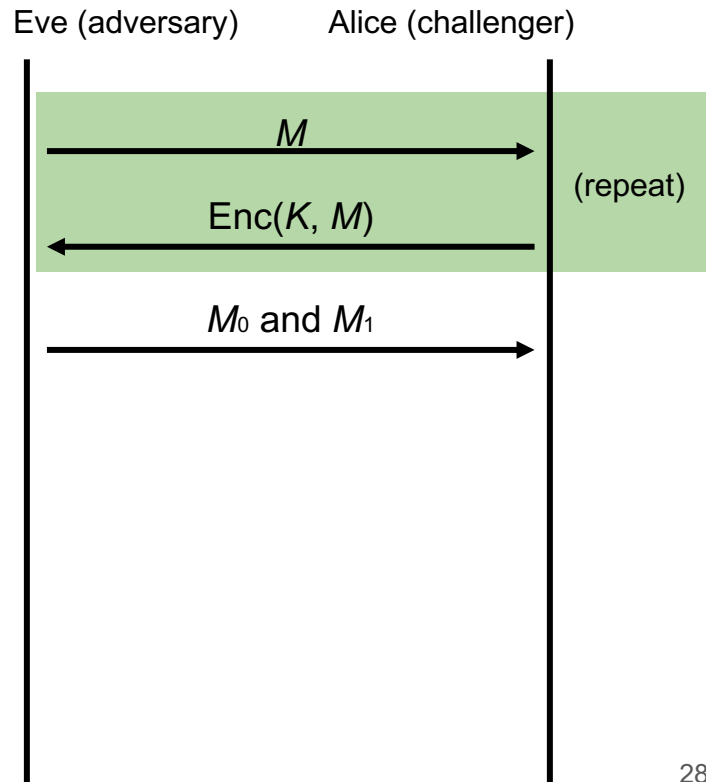
1. Eve may choose plaintexts to send to Alice and receives their ciphertexts



Defining Confidentiality: IND-CPA

ITIS 6200 / 8200

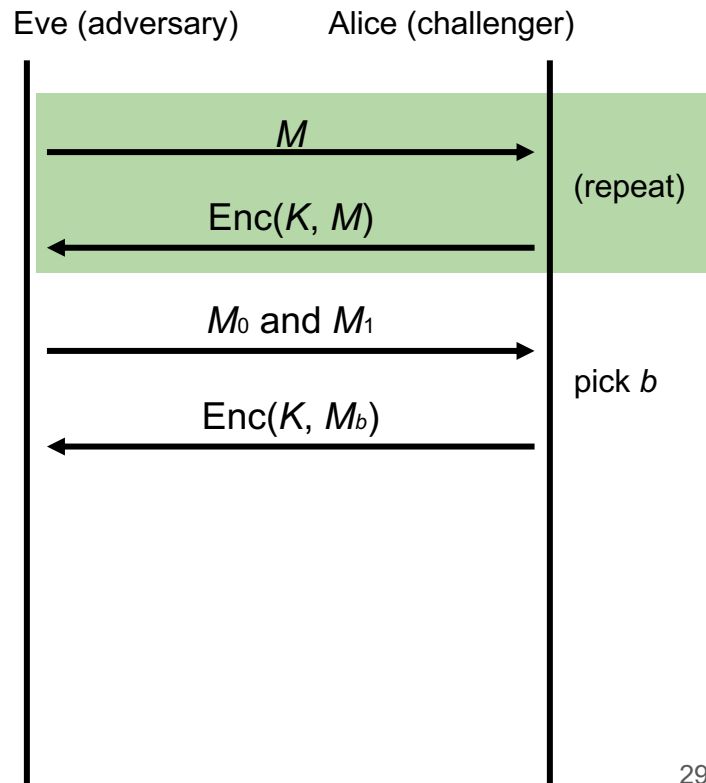
1. Eve may choose plaintexts to send to Alice and receives their ciphertexts
2. Eve issues a pair of plaintexts M_0 and M_1 to Alice



Defining Confidentiality: IND-CPA

ITIS 6200 / 8200

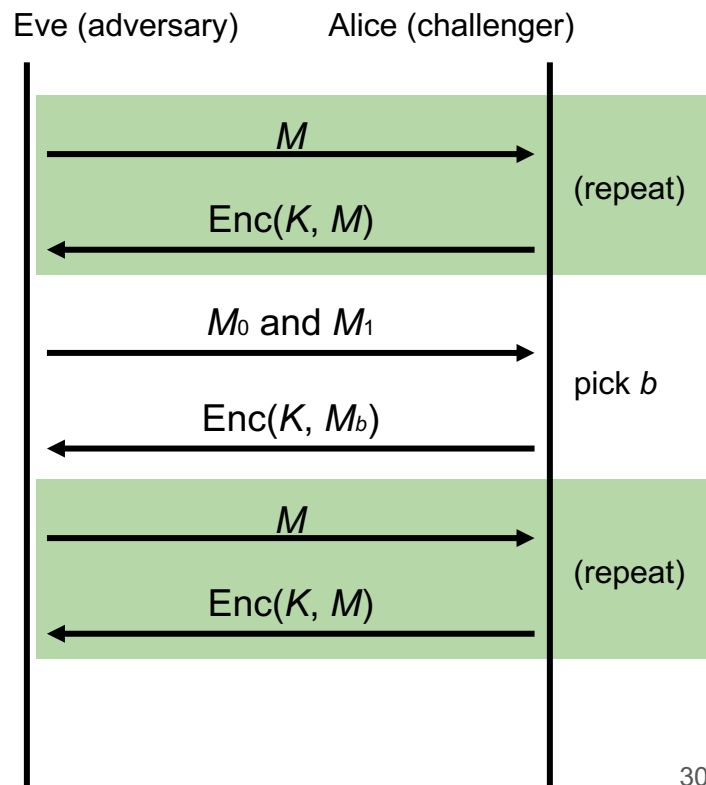
1. Eve may choose plaintexts to send to Alice and receives their ciphertexts
2. Eve issues a pair of plaintexts M_0 and M_1 to Alice
3. Alice randomly chooses either M_0 or M_1 to encrypt and sends the encryption back
 - a. Alice does not tell Eve which one was encrypted!



Defining Confidentiality: IND-CPA

ITIS 6200 / 8200

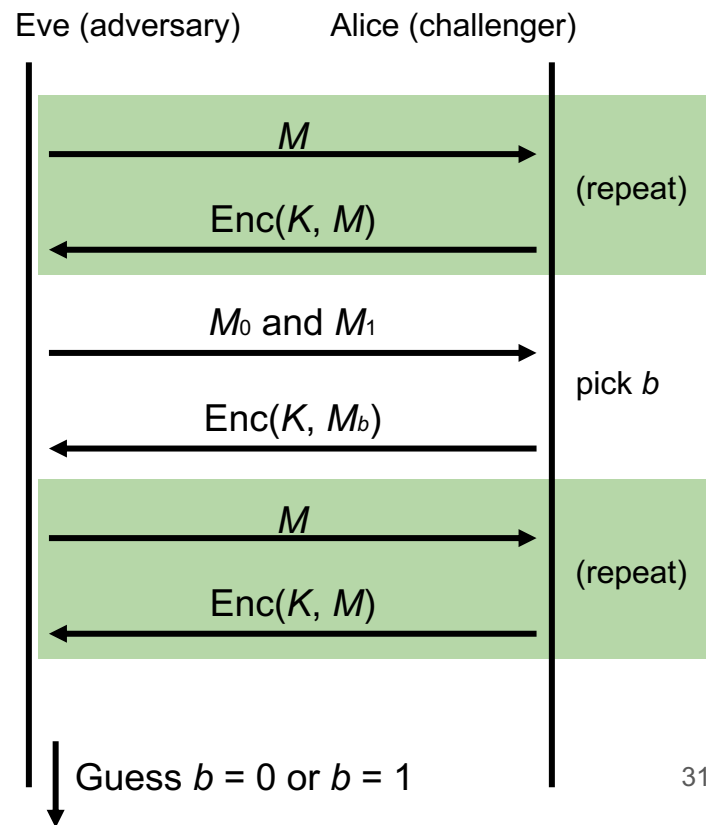
1. Eve may choose plaintexts to send to Alice and receives their ciphertexts
2. Eve issues a pair of plaintexts M_0 and M_1 to Alice
3. Alice randomly chooses either M_0 or M_1 to encrypt and sends the encryption back
 - a. Alice does not tell Eve which one was encrypted!
4. Eve may again choose plaintexts to send to Alice and receives their ciphertexts



Defining Confidentiality: IND-CPA

ITIS 6200 / 8200

1. Eve may choose plaintexts to send to Alice and receives their ciphertexts
2. Eve issues a pair of plaintexts M_0 and M_1 to Alice
3. Alice randomly chooses either M_0 or M_1 to encrypt and sends the encryption back
 - a. Alice does not tell Eve which one was encrypted!
4. Eve may again choose plaintexts to send to Alice and receives their ciphertexts
5. Eventually, Eve outputs a guess as to whether Alice encrypted M_0 or M_1



Defining Confidentiality: IND-CPA

ITIS 6200 / 8200

- If Eve correctly guesses which message Alice encrypted, then Eve wins. Otherwise, she loses.
- How does Eve guess whether M_0 or M_1 was encrypted? What strategy does she use?
 - We don't *assume* she uses a particular strategy; Eve represents all possible strategies
- Proving insecurity: There exists at least *one* strategy that can win the IND-CPA game with probability $> 1/2$
 - $1/2$ is the probability of winning by random guessing
 - If you can be better than random, then the ciphertext has leaked information, and Eve is able to learn it and use it to gain an advantage!
- Proving security: For *all* attackers/Eve-s, the probability of winning the IND-CPA game is at most $1/2$

Edge Cases: Length

ITIS 6200 / 8200

- Cryptographic schemes are (usually) allowed to leak the length of the message
 - To hide length: All messages must always be the same length
 - 16-byte messages: We can't encrypt large messages (images, videos, etc.)!
 - 1-GB messages: Sending small messages (text, Tweets, etc.) needs 1 GB of bandwidth!
 - This is impractical
- In the IND-CPA game: M_0 and M_1 must be the same length
 - To break IND-CPA, Eve must learn something other than message length



Edge Cases: Attacker Runtime

ITIS 6200 / 8200

- Some schemes are theoretically vulnerable, but secure in any real-world setting
 - If an attack takes longer than the life of the solar system to complete, it probably won't happen!
 - Or if it would require a computer made out of a literal galaxy worth of science-fiction nanotech
- In the IND-CPA game: Eve is limited to a practical runtime
 - One common practical limit: Eve is limited to polynomial runtime algorithms (no exponential-time algorithms)



Edge Cases: Negligible Advantage

ITIS 6200 / 8200

- Sometimes it's possible for Eve to win with probability $1/2 + 1/2^{128}$
 - This probability is greater than $1/2$, but it's so close to $1/2$ that it's as good as $1/2$.
 - Eve's advantage is so small that she can't use it for any practical attacks
- In the IND-CPA game: The scheme is secure even if Eve can win with probability $\leq 1/2 + \epsilon$, where ϵ is *negligible*
 - The actual mathematical definition of negligible is out of scope
 - Example: $1/2 + 1/2^{128}$: Negligible advantage
 - Example: $2/3$: Non-negligible advantage



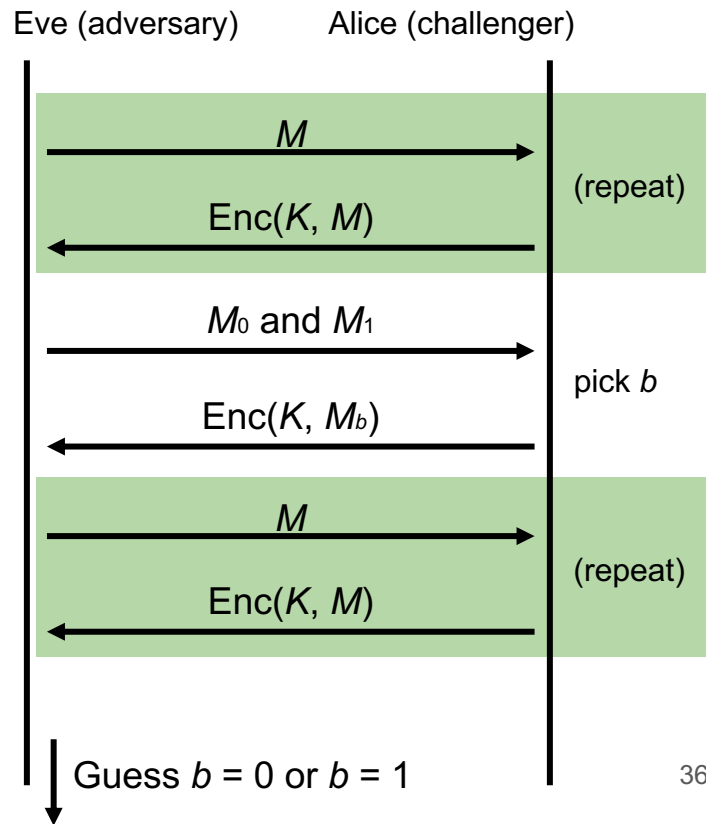
IND-CPA: Putting it together

ITIS 6200 / 8200

1. Eve may choose plaintexts to send to Alice and receives their ciphertexts
2. Eve issues a pair of plaintexts M_0 and M_1 to Alice
3. Alice randomly chooses either M_0 or M_1 to encrypt and sends the encryption back
 - Alice does not tell Eve which one was encrypted!
4. Eve may again choose plaintexts to send to Alice and receives their ciphertexts
5. Eventually, Eve outputs a guess as to whether encrypted M_0 or M_1



- An encryption scheme is IND-CPA secure if for all polynomial time attackers Eve:
 - Eve can win with probability $\leq 1/2 + \epsilon$, where ϵ is *negligible*.



A Brief History of Cryptography

Cryptography by Hand: Caesar Cipher

ITIS 6200 / 8200

- One of the earliest cryptographic schemes was the **Caesar cipher**
 - Used by Julius Caesar over 2,000 years ago
- **KeyGen()**:
 - Choose a key K randomly between 0 and 25
- **Enc(K, M)**:
 - Replace each letter in M with the letter K positions later in the alphabet
 - If $K = 3$, plaintext DOG becomes GRJ
- **Dec(K, C)**:
 - Replace each letter in C with the letter K positions earlier in the alphabet
 - If $K = 3$, ciphertext GRJ becomes DOG

$K = 3$			
M	C	M	C
A	D	N	Q
B	E	O	R
C	F	P	S
D	G	Q	T
E	H	R	U
F	I	S	V
G	J	T	W
H	K	U	X
I	L	V	Y
J	M	W	Z
K	N	X	A
L	O	Y	B
M	P	Z	C

Cryptography by Hand: Attacks on the Caesar Cipher

ITIS 6200 / 8200

- Eve sees the ciphertext JCKN ECGUCT, but doesn't know the key K
- If you were Eve, how would you try to break this algorithm?
- Brute-force attack: Try all 26 possible keys!
- Use existing knowledge: Assume that the message is in English

+1 IBJM DBFTBS

+2 HAIL CAESAR

+3 GZHK BZDRZQ

+4 FYGJ AYCQYP

+5 EXFI ZXBPXO

+6 DWEH YWAOWN

+7 CVDG XVZNVN

+8 BUCF WUYMUL

+9 ATBE VTXLTK

+10 ZSAD USWKSJ

+11 YRZC TRVJRI

+12 XQYB SQUIQH

+13 WPXA RPTHPI

+14 VOWZ QOSGOF

+15 UNVY PNRFNE

+16 TMUX OMQEMD

+17 SLTW NLPDLC

+18 RKSV MKOCKB

+19 QJRU LJNBJA

+20 PIQT KIMAIZ

+21 OHPS JHLZHY

+22 NGOR IGKYGX

+23 MFNQ HFJXFW

+24 LEMP GEIWEV

+25 KDLO FDHVDU

Cryptography by Hand: Attacks on the Caesar Cipher

ITIS 6200 / 8200

- Eve sees the ciphertext JCKN ECGUCT, but doesn't know the key K
- Chosen-plaintext attack: Eve tricks Alice into encrypting plaintext of her choice
 - Eve sends a message $M = AAA$ and receives $C = CCC$
 - Eve can deduce the key: C is 2 letters after A, so $K = 2$
 - Eve has the key, so she can decrypt the ciphertext

Cryptography by Hand: Substitution Cipher

ITIS 6200 / 8200

- A better cipher: create a mapping of each character to another character.
 - Example: A = N, B = Q, C = L, D = Z, etc.
 - Unlike the Caesar cipher, the shift is no longer constant!
- **KeyGen():**
 - Generate a random, one-to-one mapping of characters
- **Enc(K, M):**
 - Map each letter in M to the output according to the mapping K
- **Dec(K, C):**
 - Map each letter in C to the output according to the *reverse* of the mapping K

K			
M	C	M	C
A	N	N	G
B	Q	O	P
C	L	P	T
D	Z	Q	A
E	K	R	J
F	R	S	O
G	V	T	D
H	U	U	I
I	E	V	C
J	S	W	F
K	B	X	M
L	W	Y	X
M	Y	Z	H

Cryptography by Hand: Attacks on Substitution Ciphers

ITIS 6200 / 8200

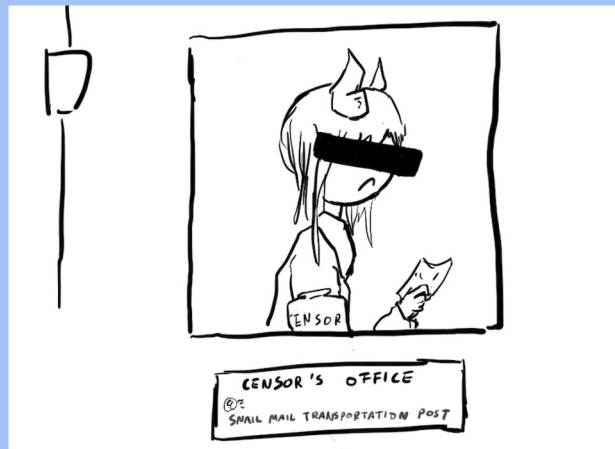
- Does the brute-force attack still work?
 - There are $26! \approx 2^{88}$ possible mappings to try
 - Too much for most modern computers... for now
- How about the chosen-plaintext attack?
 - Trick Alice into encrypting ABCDEFGHIJKLMNOPQRSTUVWXYZ, and you'll get the whole mapping!
- Another strategy: *cryptanalysis*
 - The most common English letters in text are E, T, A, O, I, N

K			
M	C	M	C
A	N	N	G
B	Q	O	P
C	L	P	T
D	Z	Q	A
E	K	R	J
F	R	S	O
G	V	T	D
H	U	U	I
I	E	V	C
J	S	W	F
K	B	X	M
L	W	Y	X
M	Y	Z	H

Takeaways

ITIS 6200 / 8200

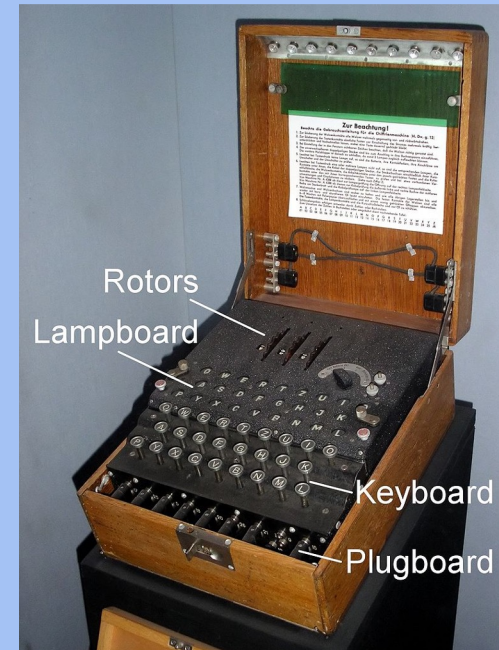
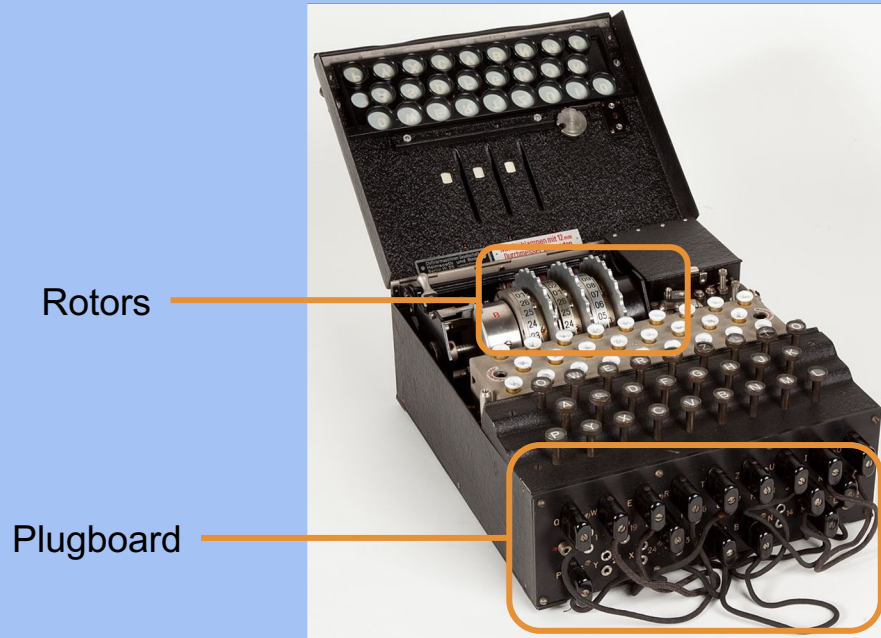
- Cryptography started with paper-and-pencil algorithms (Caesar cipher)
- Then cryptography moved to machines (Enigma)
- Finally, cryptography moved to computers (which we're about to study)
- Hopefully you gained some intuition for some of the cryptographic definitions



Cryptography by Machines: Enigma

ITIS 6200 / 8200

- A mechanical encryption machine used by the Germans in WWII



Enigma Operating Principle: Rotor Machine

ITIS 6200 / 8200

- The encryption core was composed of 3 or 4 rotors
 - Each rotor was a fixed permutation (e.g. A maps to F, B maps to Q...)
 - And the end was a "reflector", a rotor that sent things backwards
 - Plus a fixed-permutation plugboard
- A series of rotors were arranged in a sequence
 - Each keypress would generate a current from the input to one light for the output
 - Each keypress also advanced the first rotor
 - When the first rotor makes a full rotation, the second rotor advances one step
 - When the second rotor makes a full rotation, the third rotor advances once step

Cryptography by Machines: Enigma

ITIS 6200 / 8200

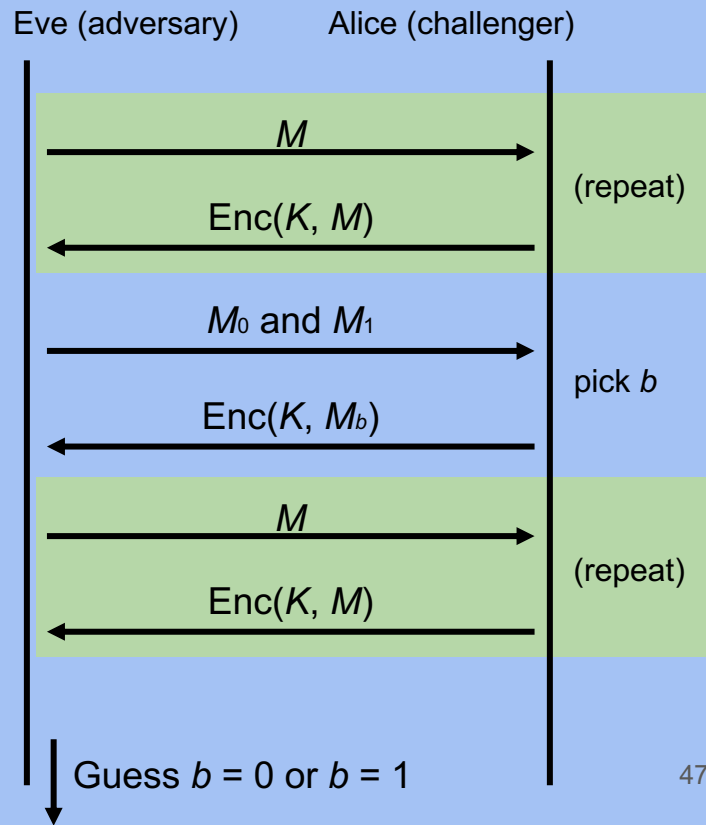
- **KeyGen():**
 - Choose rotors, rotor orders, rotor positions, and plugboard settings
 - 158,962,555,217,826,360,000 possible keys
- **Enc(K , M) and Dec(K , C):**
 - Input the rotor settings K into the Enigma machine
 - Press each letter in the input, and the lampboard will light up the corresponding output letter
 - Encryption and decryption are the same algorithm!
- Germans believed that Enigma was an “unbreakable code”



Cryptography by Machines: Enigma

ITIS 6200 / 8200

- Enigma has a significant weakness: a letter never maps to itself
 - No rotor maps a letter to itself
 - The reflector never maps a letter to itself
 - This property is necessary for Enigma's mechanical system to work
- What pair of messages should Eve send to Alice in the challenge phase?
 - Send $M_0 = A^k$, $M_1 = B^k$
 - M_0 is a string of k 'A' characters, M_1 is a string of k 'B' characters
- How can Eve probably know which message Alice encrypted?
 - If there are no 'A' characters, it was M_0
 - If there are no 'B' characters, it was M_1



Cryptography by Machines: Attack on Enigma

ITIS 6200 / 8200

- Polish and British cryptographers built BOMBE, a machine to brute-force Enigma keys
- Why was Enigma breakable?
 - Kerckhoff's principle: The Allies stole Enigma machines, so they knew the algorithm
 - Known plaintext attacks: the Germans often sent predictable messages (e.g. the weather report every morning)
 - Chosen plaintext attacks: the Allies could trick the Germans into sending a message (e.g. "newly deployed minefield")
 - Brute-force: BOMBE would try many keys until the correct one was found
 - Plus a weakness: You'd be able to try multiple keys with the same hardware configuration

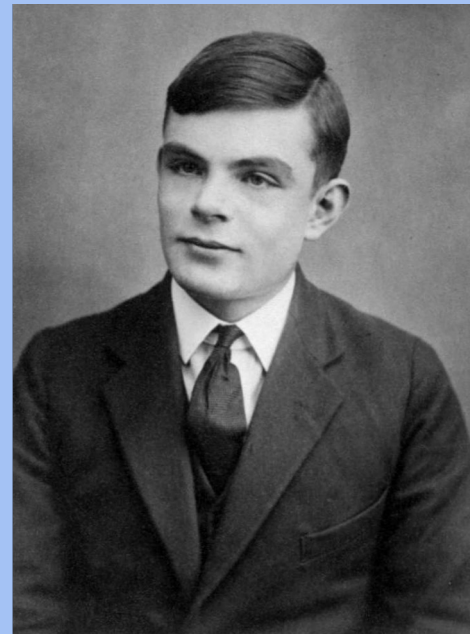


BOMBE machine

Cryptography by Machines: Legacy of Enigma

ITIS 6200 / 8200

- Alan Turing, one of the cryptographers who broke Enigma, would go on to become one of the founding fathers of computer science
- Most experts agree that the Allies breaking Enigma shortened the war in Europe by about a year



Alan Turing

Cryptography by Computers

ITIS 6200 / 8200

- The modern era of cryptography started after WWII, with the work of Claude Shannon
- “New Directions in Cryptography” (1976) showed how number theory can be used in cryptography
 - Its authors, Whitfield Diffie and Martin Hellman, won the Turing Award in 2015 for this paper
- This is the era of cryptography we’ll be focusing on



One of these is Diffie, and
the other one is Hellman.



One-Time Pads

Cryptography Roadmap

ITIS 6200 / 8200

	Symmetric-key	Asymmetric-key
Confidentiality	<ul style="list-style-type: none">● One-time pads● Block ciphers● Stream ciphers	<ul style="list-style-type: none">● RSA encryption
Integrity, Authentication	<ul style="list-style-type: none">● MACs	<ul style="list-style-type: none">● Digital signatures

- Hash functions
- Pseudorandom number generators
- Public key exchange (e.g. Diffie-Hellman)

- Key management (certificates)
- Password management

Review: XOR

ITIS 6200 / 8200

The XOR operator takes two bits and outputs one bit:

$0 \oplus 0 = 0$
$0 \oplus 1 = 1$
$1 \oplus 0 = 1$
$1 \oplus 1 = 0$

Useful properties of XOR:

$x \oplus 0 = x$
$x \oplus x = 0$
$x \oplus y = y \oplus x$
$(x \oplus y) \oplus z = x \oplus (y \oplus z)$
$(x \oplus y) \oplus x = y$

One-Time Pads: Key Generation

ITIS 6200 / 8200

Alice

K

0	1	1	0	0	1	0	1	0	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---

The key K is a randomly-chosen bitstring.

Recall: We are in the symmetric-key setting, so we'll assume Alice and Bob both know this key.

One-Time Pads: Encryption

ITIS 6200 / 8200

Alice

K	0	1	1	0	0	1	0	1	0	1	1	1
M	1	0	0	1	1	0	0	1	0	1	0	0

The plaintext M is the bitstring that Alice wants to encrypt.

Idea: Use XOR to scramble up M with the bits of K .

One-Time Pads: Encryption

ITIS 6200 / 8200

Alice

K	0	1	1	0	0	1	0	1	0	1	1	1
	\oplus	\oplus	\oplus	\oplus	\oplus	\oplus	\oplus	\oplus	\oplus	\oplus	\oplus	\oplus
M	1	0	0	1	1	0	0	1	0	1	0	0
	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓
C	1	1	1	1	1	1	0	0	0	0	1	1

Encryption algorithm: XOR each bit of K with the matching bit in M .

The ciphertext C is the encrypted bitstring that Alice sends to Bob over the insecure channel.

One-Time Pads: Decryption

ITIS 6200 / 8200

Bob

K	0	1	1	0	0	1	0	1	0	1	1	1
C	1	1	1	1	1	1	0	0	0	0	1	1

Bob receives the ciphertext C . Bob knows the key K . How does Bob recover M ?

One-Time Pads: Decryption

ITIS 6200 / 8200

Bob

K	0	1	1	0	0	1	0	1	0	1	1	1
	\oplus	\oplus	\oplus	\oplus	\oplus	\oplus	\oplus	\oplus	\oplus	\oplus	\oplus	\oplus
C	1	1	1	1	1	1	0	0	0	0	1	1
	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓
M	1	0	0	1	1	0	0	1	0	1	0	0

Decryption algorithm: XOR each bit of K with the matching bit in C .

One-Time Pad

ITIS 6200 / 8200

- KeyGen()
 - Randomly generate an n -bit key, where n is the length of your message
 - Recall: For today, we assume that Alice and Bob can securely share this key
 - For one-time pads, we generate a *new* key for every message
- $\text{Enc}(K, M) = K \oplus M$
 - Bitwise XOR M and K to produce C
 - In other words: XOR the i th bit of the plaintext with the i th bit of the key.
 - $C_i = K_i \oplus M_i$
 - Alice and Bob use a different key for each encryption (this is the “one-time” in one-time pad).
- $\text{Dec}(K, C) = K \oplus C$
 - Bitwise XOR C and K to produce M
 - $M_i = K_i \oplus C_i$

One-Time Pad: Correctness

ITIS 6200 / 8200

- Correctness: If we encrypt and then decrypt, we should get the original message back

$$\text{Enc}(K, M) = K \oplus M$$

Definition of encryption

$$\text{Dec}(K, \text{Enc}(K, M)) = \text{Dec}(K, K \oplus M)$$

Decrypting the ciphertext

$$= K \oplus (K \oplus M)$$

Definition of decryption

$$= M$$

XOR property

One-Time Pad: Security

ITIS 6200 / 8200

- Recall our definition of confidentiality: The ciphertext should not give the attacker any additional information about the plaintext
- Recall our experiment to test confidentiality from earlier:
 - Alice has encrypted and sent either M_0 or M_1
 - Eve knows either M_0 or M_1 was sent, but doesn't know which
 - Eve reads the ciphertext and tries to guess which message was sent
 - If the probability that Eve correctly guesses which message was sent is $1/2$, then the encryption scheme is confidential

One-Time Pad: Security

ITIS 6200 / 8200

Possibility 0: Alice sends $\text{Enc}(K, M_0)$

The ciphertext is $C = K \oplus M_0$

Therefore, $K = C \oplus M_0$

Possibility 1: Alice sends $\text{Enc}(K, M_1)$

The ciphertext is $C = K \oplus M_1$

Therefore, $K = C \oplus M_1$

K was chosen randomly, so both possibilities are equally possible

Eve has learned no new information, so the scheme is *perfectly secure*

Impracticality of One-Time Pads

ITIS 6200 / 8200

- Problem #1: Key generation
 - For security to hold, keys must be randomly generated for every message, and never reused
 - Randomness is expensive, as we'll see later
- Problem #2: Key distribution
 - To communicate an n -bit message, we need to securely communicate an n -bit key first
 - But if we have a way to securely communicate an n -bit key, we could have communicated the message directly!
- Only practical application: Communicate keys in advance
 - You have a secure channel now, but you won't have it later
 - Use the secure channel now to communicate keys in advance
 - Use one-time pad later to communicate over the insecure channel
 - And people can compute this by hand without computers!

Summary

ITIS 6200 / 8200

- What is cryptography?
- Definitions of key properties
- A better definition of confidentiality: IND-CPA security
- A brief history of cryptography
- Symmetric-key crypto