

Transmission Control Protocol (TCP)

Today: Transport Layer Protocols

ITIS 6200 / 8200

- Transmission Control Protocol (TCP): Reliably sending packets

Review: IP Reliability

ITIS 6200 / 8200

- **Reliability** ensures that packets are received correctly
- IP is **unreliable** and only provides a **best effort** delivery service, which means:
 - Packets may be lost (“dropped”)
 - Packets may be corrupted
 - Packets may be delivered out of order
- It is up to higher level protocols to ensure that the connection is reliable

Scratchpad: Let's Design It Together

ITIS 6200 / 8200

- Problem: IP packets have a limited size. To send longer messages, we have to manually break messages into packets
 - When sending packets: TCP will automatically split up messages
 - When receiving packets: TCP will automatically reassemble the packets
- Problem: Packets can arrive out of order
 - When sending packets: TCP labels each byte of the message with increasing numbers
 - When receiving packets: TCP can use the numbers to rearrange bytes in the correct order
- Problem: Packets can be dropped
 - When receiving packets: TCP sends an extra message acknowledging that a packet has been received
 - When sending packets: If the acknowledgement doesn't arrive, re-send the packet

Transmission Control Protocol (TCP)

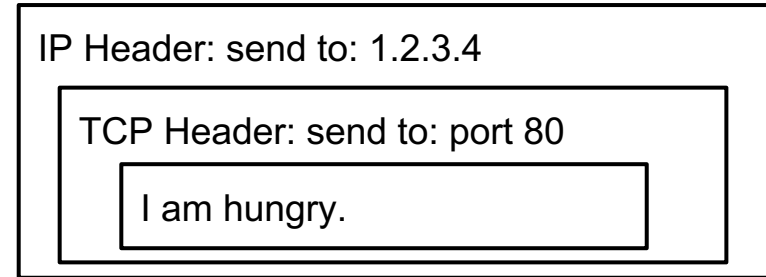
ITIS 6200 / 8200

- Provides a byte stream abstraction
 - Bytes go in one end of the stream at the source and come out at the other end at the destination
 - TCP automatically breaks streams into **segments**, which are sent as layer 3 packets
- Provides ordering
 - Segments contain sequence numbers, so the destination can reassemble the stream in order
- Provides reliability
 - The destination sends **acknowledgements** (ACKs) for each sequence number received
 - If the source doesn't receive the ACK, the source sends the packet again
- Provides ports
 - Multiple services can share the same IP address by using different ports

Ports

ITIS 6200 / 8200

- **Ports** help us distinguish between different applications on the same computer or server
 - On private computers, port numbers can be random
 - On public servers, port numbers should be constant and well-known (so users can access the right port)
- Remember: TCP is built on top of IP, so the IP header (and therefore the IP address) is still present



Establishing Sequence Numbers

ITIS 6200 / 8200

- Each TCP connection requires two sets of sequence numbers
 - One sequence number for messages from the client to the server
 - One sequence number for messages from the server to the client
- Before starting a TCP connection, the client and server must agree on two **initial sequence numbers (ISNs)**
 - The ISNs are different and random for every connection (for security reasons)

H	e	l	l	o		s	e	r	v	e	r
50	51	52	53	54	55	56	57	58	59	60	61

Messages from the client are numbered starting at 50.


H	e	l	l	o		c	l	i	e	n	t
25	26	27	28	29	30	31	32	33	34	35	36

Messages from the server are numbered starting at 25.

TCP: 3-Way Handshake

ITIS 6200 / 8200

1. Client chooses an initial sequence number x its bytes and sends a SYN (synchronize) packet to the server
2. Server chooses an initial sequence number y for its bytes and responds with a SYN-ACK packet
3. Client then returns with an ACK packet
4. Once both hosts have synchronized sequence numbers, the connection is “established”



TCP: Sending and Receiving Data

ITIS 6200 / 8200

- The TCP handlers on each side track which TCP segments have been received for each connection
 - A connection is identified by these 5 values (sometimes called a 5-tuple)
 - Source IP
 - Destination IP
 - Source Port
 - Destination Port
 - Protocol
- Data from the bytestream can be presented to the application when all data before has been received and presented
 - Recall: TCP presents data to the application as a bytestream, so the order must be preserved from one end to the other, even if packets are received out of order

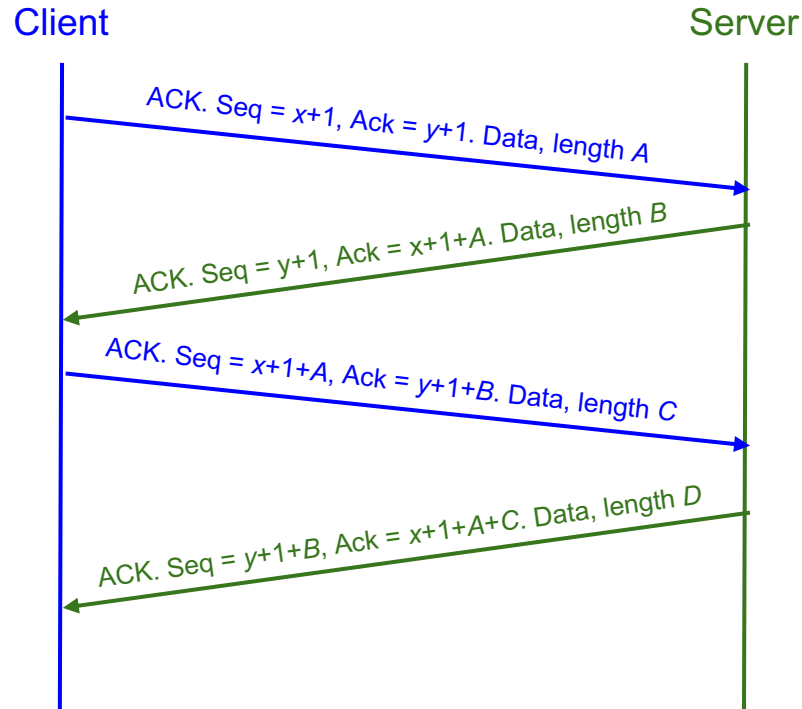
TCP: Sending and Receiving Data

ITIS 6200 / 8200

- Byte i of the bytestream is represented by sequence number $x + i$
 - The first byte is byte $i = 1$, since sequence number x was used for the SYN packet and y for the SYN-ACK packet
- A packet's sequence number is the number of the first byte of its data
 - This number is from the sender's set of sequence numbers
- A packet's ACK number, if the ACK flag is set, is the number of the byte immediately after the last received byte
 - This number is from the receiver's set of sequence numbers
 - This would be (sequence number) + (length of data) for the last received packet

TCP: Sending and Receiving Data

ITIS 6200 / 8200



TCP: Retransmission

ITIS 6200 / 8200

- If a packet is dropped (lost in transit):
 - The recipient will not send an ACK, so the sender will not receive the ACK
 - The sender repeatedly tries to send the packet again until it receives the ACK
- If a packet is received, but the ACK is dropped:
 - The sender tries to send the packet again since it didn't receive the ACK
 - The recipient ignores the duplicate data and sends the ACK again
- When packets are dropped in TCP, TCP assumes that there is congestion and sends the data at a slower rate

TCP: Ending/Aborting a Connection

ITIS 6200 / 8200

- To **end** a connection, one side sends a packet with the FIN (finish) flag set, which should then be acknowledged
 - This means “I will no longer be sending any more packets, but I will continue to receive packets”
 - Once the other side is no longer sending packets, it sends a packet with the FIN flag set
- To **abort** a connection, one side sends a packet with the RST (reset) flag set
 - This means “I will no longer be sending nor receiving packets on this connection”
 - RST packets are not acknowledged since they usually mean that something went wrong

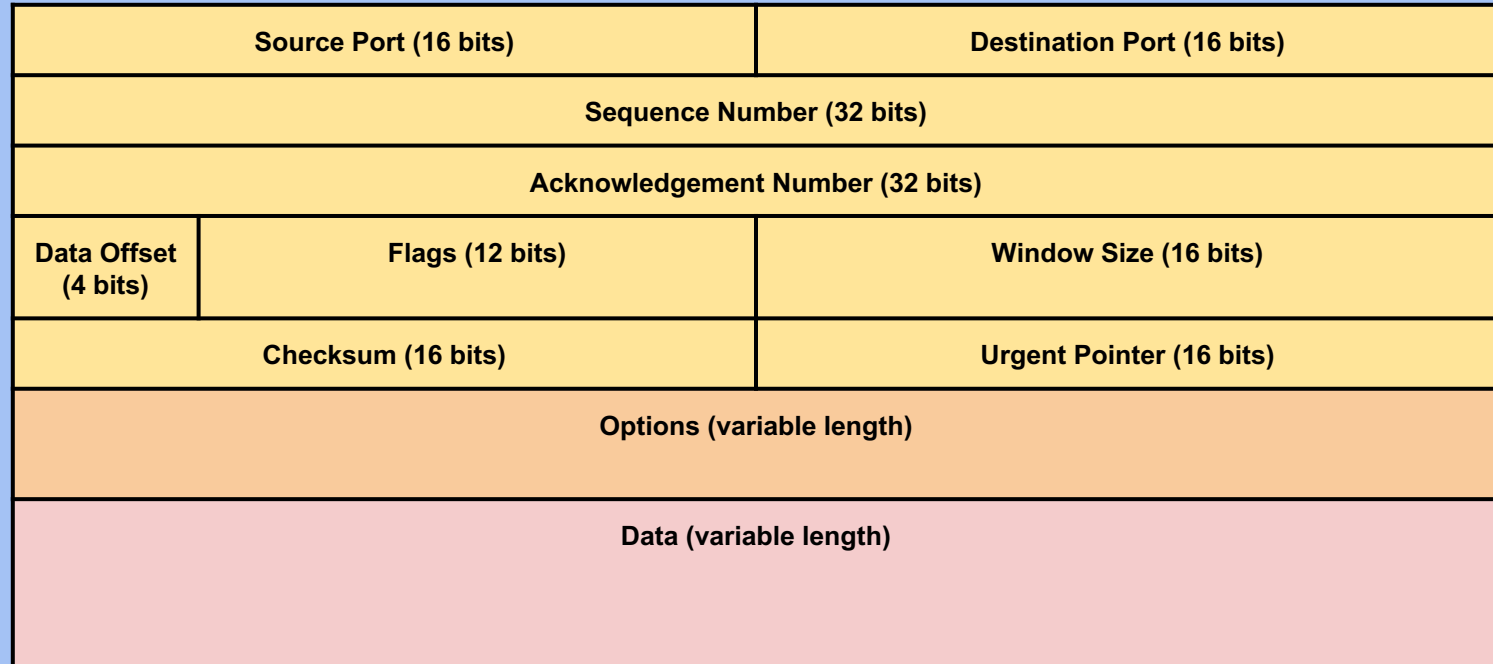
TCP Flags

ITIS 6200 / 8200

- **ACK**
 - Indicator that the user is acknowledging the receipt of something (in the ack number)
 - Pretty much always set except the very first packet
- **SYN**
 - Indicator of the beginning of the connection
- **FIN**
 - One way to end the connection
 - Requires an acknowledgement
 - No longer sending packets, but will continue to receive
- **RST**
 - One way to end a connection
 - Does not require an acknowledgement
 - No longer sending or receiving packets

TCP Segment Structure

ITIS 6200 / 8200



TCP segment header

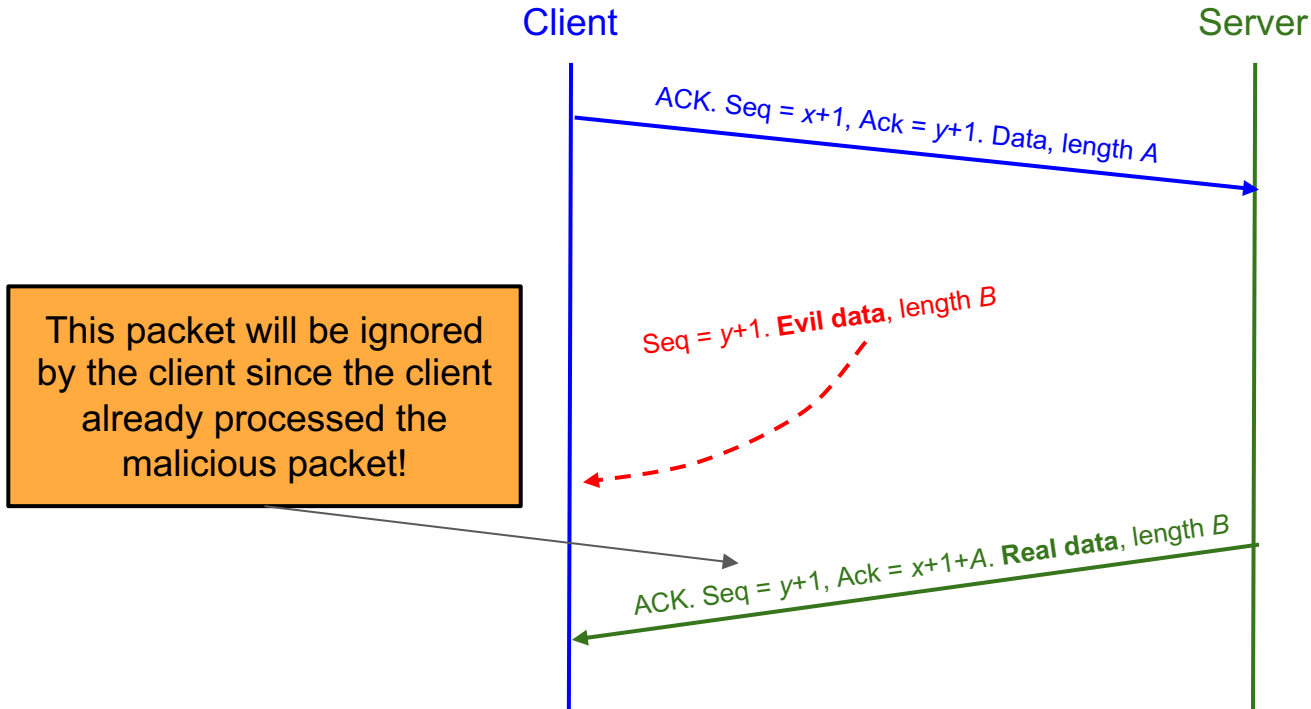
TCP Attacks

ITIS 6200 / 8200

- **TCP hijacking:** Tampering with an existing session to modify or inject data into a connection
 - **Data injection:** Spoofing packets to inject malicious data into a connection
 - Need to know: The sender's sequence number
 - Easy for MITM and on-path attackers, but off-path attackers must guess 32-bit sequence number (called **blind injection/hijacking**, considered difficult)
 - For on-path attackers, this becomes a race condition since they must beat the server's legitimate response
 - **RST injection:** Spoofing a RST packet to forcibly terminate a connection
 - Same requirements as packet injection, so easy for on-path and MITM attackers, but hard for off-path attackers
 - Often used in censorship scenarios to block access to sites

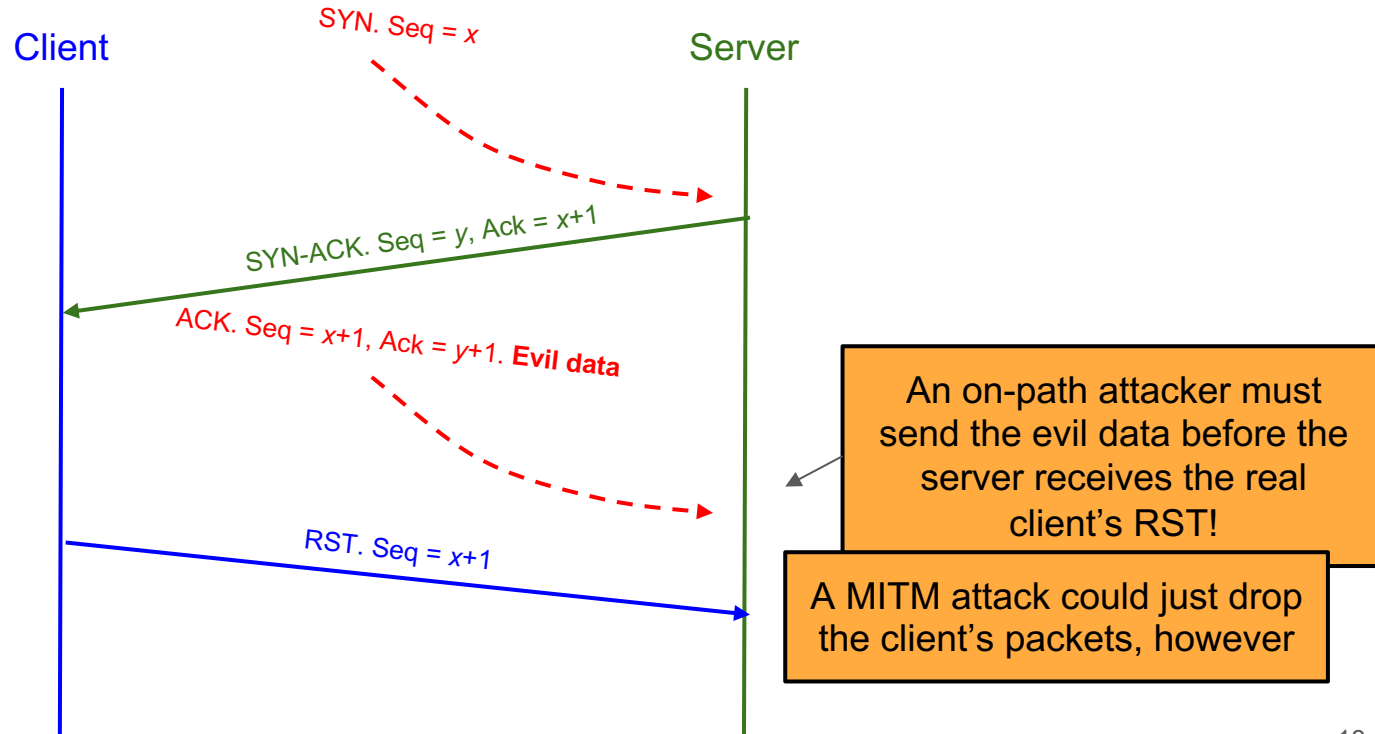
TCP Data Injection

ITIS 6200 / 8200



TCP Spoofing

ITIS 6200 / 8200



TCP Attacks

ITIS 6200 / 8200

- **TCP spoofing:** Spoofing a TCP connection to appear to come from another source IP address
 - Need to know: Sequence number in the server's response SYN-ACK packet
 - Easy for MITM and on-path attackers, but off-path attackers must guess 32-bit sequence number (called **blind spoofing**, also considered difficult)
 - For on-path attackers, this is a race condition, since the real client will send a RST upon receiving the server's SYN-ACK!

TCP Attacks

ITIS 6200 / 8200

- TCP provides no confidentiality or integrity
 - Instead, we rely on higher layers to prevent those kind of attacks
- Defense against off-path attackers rely on choosing random sequence numbers
 - Bad randomness can lead to trivial off-path attacks: TCP sequence numbers used to be based on the system clock!

Summary

ITIS 6200 / 8200

- Transmission Control Protocol (TCP): Reliably sending packets
 - 3-way handshake: Client sends SYN, server sends SYN-ACK, client sends ACK
 - Provides reliability, ordering, and ports
 - Attack: TCP hijacking through data injection or RST injection
 - Blind attacks must guess the client's or server's sequence numbers
 - Attack: TCP spoofing by sending a spoofed SYN packet
 - Blind attacks must guess the server's sequence number